

NoiseFloor — Presentation Preparation Guide

Prepared: August 19, 2026 | For: Tomorrow's Capstone Presentation

QUESTION 1: How did we perform the 4 hypothesis tests?

The Short Version (Say this in your presentation)

"We used standard statistical methods from SciPy — Python's scientific computing library — inside our PySpark notebook in Google Colab. We ran 4 tests on 148,180 filtered Nepal events."

The Detailed Breakdown

HYP 01 — "India is more negative than China" (REJECTED H_0)

- **Test Used:** Welch's Two-Sample T-Test (also called an independent samples t-test)
- **What we did:** We took ALL of India's daily average tone scores across ALL 3 crises and put them in one list. We did the same for China. Then we asked the math: "Could these two lists have come from the same population?" (i.e., are they really the same, and we just got unlucky with random variation?)
- **Result:** t-stat = -13.95, p-value = 0.0051
- **What this means:** The p-value is 0.0051, which is way below

0.05. This means there is only a 0.51% chance that India and China have the same tone. In other words, the difference is REAL, not a coincidence.

- **Plain English:** "Indian news about Nepal is scientifically, provably more negative than Chinese news. This is not random — it's structural."

HYP 02 — "India hasn't changed in 10 years" (STABLE)

- **Test Used:** Paired T-Test (comparing India 2015 vs India 2025)
- **What we did:** We took India's average tone during the 2015 crises (-2.26 for earthquake, -2.43 for blockade) and compared it to India's tone during 2025 Gen-Z protests (-2.30). We asked: "Did India's tone change over 10 years?"
- **Result:** t-stat = -0.24, p-value = 0.8094
- **What this means:** The p-value is 0.81, which is WAY above 0.05. This means we CANNOT say India changed. The difference between 2015 and 2025 is so small it could easily be random noise.
- **Plain English:** "India's negativity toward Nepal is not a reaction to specific events — it's a permanent, structural bias that hasn't budged in a decade."

HYP 03 — "India and China don't copy each other" (CONFIRMED)

- **Test Used:** Pearson Correlation Coefficient
- **What we did:** We took EACH DAY's tone for India and China during the 2025 protests (where we had day-by-day data for both). We asked: "When India's tone goes up, does China's go up too? When India dips, does China dip?"
- **Result:** $r = -0.0235$, p-value = 0.824
- **What this means:** $r = -0.0235$ is essentially ZERO. A perfect correlation would be $r = 1.0$ or $r = -1.0$. Our result of -0.02 means there is virtually NO relationship between their day-to-day coverage.
- **Plain English:** "India and China are literally watching different Nepals. When Indian news panics about something, Chinese news doesn't even notice. They cover completely different stories, different angles, different events."

HYP 04 — "China is slowly getting more negative" (REJECTED H_0)

- **Test Used:** Linear Regression (Ordinary Least Squares)
- **What we did:** We plotted China's average tone across the 3 crises in chronological order: -0.29 (2015 Earthquake) → -0.45 (2015 Blockade) → -0.71 (2025 Protests). We fitted a regression line through these points and tested if the slope is significantly different from zero.
- **Result:** t-stat = -2.90, p-value = 0.015
- **What this means:** The p-value is 0.015, which is below 0.05. The downward slope is real — China IS getting more negative over time.
- **Plain English:** "China used to be almost neutral about Nepal. But over the last 10 years, their coverage has been slowly drifting toward negativity. If this trend continues, China's media tone will eventually match India's — but they're not there yet."

If the professor asks "What library did you use?"

Answer: "We used SciPy's `ttest_ind` for the t-tests, `pearsonr` for the correlation, and `linregress` from `scipy.stats` for the regression. All computations were done in PySpark on Google Colab."

QUESTION 2: How did we get this data? How authentic is it?

Where the data comes from

- **Source:** The GDELT Project (Global Database of Events, Language, and Tone)
- **Website:** <https://www.gdeltproject.org>
- **What it is:** GDELT is one of the largest open-access databases in the world. It monitors news media from virtually every country on

earth in over 100 languages. It has been running since 2015 and processes millions of news articles per day.

- **Who runs it:** It is a research project supported by Google and maintained by Georgetown University and various research institutions. It is widely used in academic research, political science, and journalism.
- **Is it authentic?** YES. GDELT is used by the United Nations, World Bank, academic researchers worldwide, and major news organizations. It is one of the most credible open data sources for media analysis in existence.

How we accessed it

1. GDELT publishes its data as Parquet files on Amazon Web Services (AWS S3) — a public cloud storage bucket.
2. We wrote a PySpark script in Google Colab that downloaded 22,000 raw Parquet files directly from the GDELT S3 bucket.
3. Each Parquet file contains millions of "events" — each event is one news article about one interaction between two countries.
4. Total raw records downloaded: 35,528,488 events across 10 years (2015-2025).

How we filtered it

1. **Step 1 — Nepal Filter:** We filtered for events where Actor1Geo_CountryCode = 'NPL' (events happening in Nepal). This reduced 35.5 million records down to 148,180 Nepal-specific events.
2. **Step 2 — Country Split:** We further split these 148,180 events by Actor1CountryCode to separate India's coverage (IND) from China's coverage (CHN) from 94 other countries.
3. **Step 3 — Crisis Windows:** We filtered by date ranges to isolate the 3 specific crises:
 - 2015 Earthquake: April-June 2015
 - 2015 Blockade: September-December 2015
 - 2025 Gen-Z Protests: September-November 2025
4. **Step 4 — Tone Extraction:** For each event, GDELT provides a pre-computed AvgTone score (ranges from about -10 to +10). We

- averaged these scores by day to get daily tone trends.
5. **Step 5 — Output:** We exported 8 clean CSV files from this pipeline, which feed the dashboard.

If the professor asks "Did you verify the data?"

Answer: "GDELT's tone scores are computed using established NLP algorithms applied to the full text of each news article. The GDELT project has been peer-reviewed and is cited in over 4,000 academic papers. We trust the source, and our statistical tests confirm the patterns are real and not random noise."

QUESTION 3: What analyses did we perform?

In the Jupyter Notebook (NoiseFloor.ipynb)

Analysis	What It Does
PySpark ETL Pipeline	Downloads 22,000 Parquet files from AWS, filters for Nepal, extracts India/China events
Daily Sentiment Aggregation	Calculates average tone and event count per day, computes 30-day rolling averages
Spike Detection	Identifies days where event volume exceeds 3x the rolling average (found the Sept 9 spike at 30.99x)
Three-Way Incident Comparison	Computes India vs China tone across all 3 crises side by side
T0 Normalization	Aligns all 3 crises to Day 0 (the start of each crisis) so they can be compared on the same timeline

96-Country Sentiment Analysis	Ranks all 96 countries that covered Nepal by their average tone (discovered Bhutan is the only positive one at +1.64)
CAMEO QuadClass Analysis	Categorizes all events into 4 types: Verbal Cooperation, Material Cooperation, Verbal Conflict, Material Conflict
Welch's T-Test	Tests if India and China tones are statistically different
Paired T-Test (Stability)	Tests if India's tone changed between 2015 and 2025
Pearson Correlation	Tests if India and China's day-to-day coverage is correlated
Linear Regression (Drift)	Tests if China's tone is trending downward over time

In the Web Dashboard (gdelt-dashboard/)

Feature	What It Does
Main Dashboard	Shows all KPIs, 3 crisis cards, 4 hypothesis result cards, and the 10-year time series chart
Earthquake 2015 Page	Shows India vs China tone timeline during the earthquake, with detailed explanations
Blockade 2015 Page	Shows a bar chart of QuadClass categories (cooperation vs conflict) and India vs China tone timeline
Gen-Z 2025 Page	Shows India vs China daily tone waveform during Sept 2025, with the correlation finding
Hypothesis Engine	Interactive tool: select a crisis, click RUN TEST, and watch the density distribution curves shift
Data Sources	Shows the data pipeline architecture, top 10 source countries chart, and geographic reach

Methodology	Beginner-friendly explanation of how we measured tone, filtered data, and what P-values mean
Export/Data Explorer	Searchable table of all raw data with dataset switching

QUESTION 4: Which files to copy, how to run, and how to stop

Files to Copy (The MINIMUM you need)

Copy the ENTIRE gdelt-dashboard folder. That's it. Here's what's inside:

```
gdelt-dashboard/  
├─ index.html           ← The main webpage (all HTML)  
├─ style.css            ← All the styling (colors, layout)  
├─ app.js               ← All the JavaScript (charts, int  
├─ data_bundle.js       ← All your data (compiled from th  
└─ data.js              ← Old data file (not used, but ha
```

Total size: ~250 KB. That's it. No install, no dependencies, no npm.

ALSO keep these for reference (not needed to run the dashboard):

- NoiseFloor.ipynb — Your Jupyter notebook (the backend work)
- NoiseFloor.md — The original reference guide
- The 8 CSV files — Your raw data outputs
- PRESENTATION_PREP.md — This file you're reading now

How to RUN the project

Option A: VS Code Live Server (EASIEST for presentation)

1. Open VS Code
2. Install the "Live Server" extension (by Ritwick Dey) from Extensions panel
3. Open the gdelt-dashboard folder in VS Code
4. Right-click index.html → "Open with Live Server"
5. Browser opens automatically at <http://127.0.0.1:5500>

Option B: Python HTTP Server (No extensions needed)

1. Open Terminal / Command Prompt / PowerShell
2. Navigate to the gdelt-dashboard folder: `cd C:\Users\user\Downloads\BIGGGG\gdelt-dashboard`
3. Run: `python -m http.server 8000`
4. Open browser and go to: <http://localhost:8000>

Option C: Node.js (If you have Node installed)

1. Open Terminal in the gdelt-dashboard folder
2. Run: `npx serve .`
3. It will tell you the URL to open

How to STOP the project

If using Live Server:

- Click the "Port: 5500" button in VS Code's bottom status bar, OR
- Press `Ctrl+Shift+P` → type "Live Server: Stop" → Enter

If using Python server:

- Go to the terminal where it's running
- Press `Ctrl + C`
- That's it, it stops immediately

If using Node (`npx serve`):

- Go to the terminal where it's running
- Press Ctrl + C

IMPORTANT WARNING:

You CANNOT just double-click index.html to open it directly. Browsers block loading of local .js files for security. You MUST use a server (Live Server, Python, or Node). If you double-click, you'll see a blank page.

QUESTION 5: What can be added in 2 hours?

After reviewing NoiseFloor.md and the current state of the dashboard, here are the most realistic and impactful additions:

TOP 3 RECOMMENDED (In order of priority)

1. China Drift Mini-Chart (30 minutes)

Add a simple visual showing China's tone drifting downward across the 3 crises:

- Earthquake 2015: -0.29
- Blockade 2015: -0.45
- Gen-Z 2025: -0.71

This would be a small stepped line chart or 3 colored bars, added either to the Dashboard or the Hypothesis Engine page. It directly proves HYP-04 visually and is very easy to explain in a presentation.

Why it matters: Right now the China Drift finding is only mentioned in text. A visual makes it 10x more convincing.

2. Top 10 Divergence Days Table (30 minutes)

Add a card showing the 10 days where India and China's tones were MOST different. For example: "On Oct 15, 2025: India tone was -8.2, China tone was +1.3 → Gap of 9.5 points"

This data already exists in `india_china_comparison.csv`. Just sort by the absolute difference and display the top 10.

Why it matters: It turns the abstract correlation finding ($r = -0.02$) into a concrete, jaw-dropping list of specific days.

3. Bhutan Anomaly Highlight Card (20 minutes)

Add a callout card to the Dashboard or Data Sources page: "Out of 96 countries that covered Nepal, only ONE has a positive average tone: Bhutan (+1.64). Every other country, including so-called allies, covers Nepal negatively."

This data is already in `source_countries.csv`.

Why it matters: This is a dramatic, memorable finding that will stick in the professor's mind.

TIME BUDGET:

- China Drift Chart: 30 min
- Top 10 Divergence Table: 30 min
- Bhutan Anomaly Card: 20 min
- Testing & Polish: 40 min
- TOTAL: 2 hours

NOT recommended for 2 hours:

- World Map / Choropleth — Too complex, needs D3.js/Leaflet setup (full day minimum)
- ML Sentiment Model — Needs training data and model iteration (half day minimum)
- NLP Reanalysis — Requires API setup and significant coding (full day)

- Network Graph — D3 force-directed graphs are finicky to get right (half day)
-

BONUS: Demo Script for Presentation

Here's a suggested flow for your live demo:

1. START ON DASHBOARD (2 min): "This is our command center. We processed 35.5 million news events. Here you can see India's tone is -2.33, China is -0.48. The gap is confirmed real with $p = 0.005$."
2. CLICK EARTHQUAKE 2015 (2 min): "During the earthquake, both countries sent aid. But look at the two lines — India (red) stays much lower than China (cyan). Indian media was criticized for being PR-driven, which sparked #GoHomeIndianMedia."
3. CLICK BLOCKADE 2015 (2 min): "This was the most hostile event. Indian media downplayed it, Chinese media called it hegemony. Look at the Material Conflict bar — it spiked massively."
4. CLICK GEN-Z 2025 (2 min): "This is the most recent crisis. Notice how the red and cyan lines move completely independently — that's the $r = -0.02$ correlation. They're literally watching different Nepals."
5. CLICK HYPOTHESIS ENGINE (2 min): "Here we can test each crisis individually. Watch what happens when I select 'Blockade' and hit RUN TEST..." *click* "...the curves shift to show the blockade-specific distribution."
6. CLICK METHODOLOGY (1 min): "This page explains our entire pipeline in plain English."

Total demo time: ~11 minutes. Perfect for a 15-minute slot with Q&A.

FINAL CHECKLIST BEFORE TOMORROW

☐ Copy gdelt-dashboard folder to a USB drive as backup ☐ Test that Live Server / Python server works on the presentation laptop ☐ Practice the demo flow at least once ☐ Make sure you can explain what a T-Test is in one sentence: "It tells us if the difference between two groups is real or just random luck" ☐ Make sure you can explain what a P-value is: "It's the probability that we got this result by pure chance. Below 0.05 = it's real." ☐ Make sure you can explain what Pearson Correlation is: "It measures if two things move together. 1.0 = perfectly linked, 0 = no relationship at all." ☐ Make sure the RUN TEST button works in Hypothesis Engine ☐ Make sure all 8 sidebar tabs load without errors